

life-insurance-churn-prediction-1

November 25, 2023

```
[1]: !pip install keras
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.utils import resample
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
```

Requirement already satisfied: keras in c:\users\neha\anaconda3\lib\site-packages (2.13.1)

1 DATA CLEANING AND PREPROCESSING

```
[2]: data = pd.read_csv('randomdata.csv')
```

2 #Data Exploration

```
[3]: data.head()
```

```
[3]: Unnamed: 0    Customer Name \
0          0  Christine Payne
1          1   Tony Fernandez
2          2  Christopher Kim
3          3    Nicole Allen
4          4    Linda Cruz

                                Customer_Address \
0  7627 Anderson Rest Apt. 265,Lake Heather, DC 3...
1  3953 Cindy Brook Apt. 147,East Lindatown, TN 4...
2      8693 Walters Mountains,South Tony, TX 88407
3      56926 Webster Coves,Shawnmouth, NV 04853
4      489 Thomas Forges Apt. 305,Jesseton, GA 36765

                                Company Name Claim Reason Data confidentiality \
```

0	Williams, Henderson and Perez	Travel	Low
1	Moore-Goodwin	Medical	High
2	Smith-Holmes	Phone	Medium
3	Harrell-Perez	Phone	Medium
4	Simpson, Kramer and Hughes	Phone	Medium

	Claim Amount	Category Premium	Premium/Amount Ratio	Claim Request output \
0	377	4794	0.078640	No
1	1440	14390	0.100069	No
2	256	1875	0.136533	No
3	233	1875	0.124267	No
4	239	1875	0.127467	No

	BMI	Churn
0	21	Yes
1	24	Yes
2	18	Yes
3	24	Yes
4	21	Yes

```
[4]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200000 entries, 0 to 199999
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Unnamed: 0            200000 non-null int64
1   Customer Name         200000 non-null object
2   Customer_Address     200000 non-null object
3   Company Name         200000 non-null object
4   Claim Reason         200000 non-null object
5   Data confidentiality  200000 non-null object
6   Claim Amount         200000 non-null int64
7   Category Premium     200000 non-null int64
8   Premium/Amount Ratio  200000 non-null float64
9   Claim Request output  200000 non-null object
10  BMI                   200000 non-null int64
11  Churn                 200000 non-null object
dtypes: float64(1), int64(4), object(7)
memory usage: 18.3+ MB
```

```
[5]: data.describe()
```

```
[5]:
```

	Unnamed: 0	Claim Amount	Category Premium	Premium/Amount Ratio \
count	200000.000000	200000.000000	200000.000000	200000.000000
mean	99999.500000	1120.478840	8963.783895	0.125024

std	57735.171256	796.660796	6114.737202	0.034742
min	0.000000	1.000000	399.000000	0.002506
25%	49999.750000	245.000000	1875.000000	0.106741
50%	99999.500000	1390.000000	14390.000000	0.125122
75%	149999.250000	1844.000000	14390.000000	0.143155
max	199999.000000	2299.000000	14390.000000	0.248120

	BMI
count	200000.000000
mean	23.007205
std	3.164976
min	18.000000
25%	20.000000
50%	23.000000
75%	26.000000
max	28.000000

3 # Handling Missing Data

```
[6]: # Step 1: Remove Duplicate Rows
data.drop_duplicates(inplace=True)
```

```
[7]: # Step 2: Remove Irrelevant Columns
# Identify and drop columns that are not relevant for churn prediction
irrelevant_columns = ['Customer Name', 'Customer_Address', 'Company Name',
↳ 'Data confidentiality', 'Claim Amount', 'Category Premium', 'Premium/Amount',
↳ 'Ratio']
data.drop(columns=irrelevant_columns, inplace=True)
```

```
[8]: # Step 3: Data Preprocessing
# After removing duplicates and irrelevant columns, you may proceed with data
↳ preprocessing
# This may include handling missing data, encoding categorical variables, and
↳ scaling/normalizing numerical features

# Handling Missing Data
data.dropna(subset=['Churn'], inplace=True)
```

```
[9]: # Encoding Categorical Variables
# Identify categorical columns in your dataset
import pandas as pd
from sklearn.preprocessing import LabelEncoder
# Load your dataset
data = pd.read_csv('randomdata.csv')
# Identify categorical columns in your dataset
categorical_columns = ['Company Name', 'Claim Reason', 'Category Premium']
```

```
# Perform label encoding for categorical columns
for column in categorical_columns:
    le = LabelEncoder()
    data[column] = le.fit_transform(data[column])
# Display the resulting DataFrame
print(data)
```

```

      Unnamed: 0    Customer Name \
0              0  Christine Payne
1              1   Tony Fernandez
2              2  Christopher Kim
3              3   Nicole Allen
4              4    Linda Cruz
...
199995      199995  Matthew Estrada
199996      199996      James Bean
199997      199997      David Meyer
199998      199998      Martha Stone
199999      199999      Shannon Lewis
```

```

      Customer_Address  Company Name \
0  7627 Anderson Rest Apt. 265,Lake Heather, DC 3...  122584
1  3953 Cindy Brook Apt. 147,East Lindatown, TN 4...  77347
2      8693 Walters Mountains,South Tony, TX 88407    106968
3      56926 Webster Coves,Shawnmouth, NV 04853    44952
4      489 Thomas Forges Apt. 305,Jesseton, GA 36765  104639
...
199995      2024 Lopez Gateway,Lake Pamelaafort, MS 35772    16350
199996      0268 Lori Falls,West Jeffrey, SC 49142    114158
199997      00573 Miller Cliff,New Allenbury, SC 68902    104547
199998      62681 Peters Cove,South Anthony, RI 99783    4850
199999      Unit 6569 Box 2236,DPO AE 88045    98869
```

```

      Claim Reason Data confidentiality  Claim Amount  Category Premium \
0              3              Low          377          2
1              0              High         1440          3
2              2             Medium          256          1
3              2             Medium          233          1
4              2             Medium          239          1
...
199995      0              High         1563          3
199996      0              High         1342          3
199997      0              High         2278          3
199998      3              Low          532          2
199999      0              High         1755          3
```

Premium/Amount Ratio Claim Request output BMI Churn

0	0.078640	No	21	Yes
1	0.100069	No	24	Yes
2	0.136533	No	18	Yes
3	0.124267	No	24	Yes
4	0.127467	No	21	Yes
...
199995	0.108617	No	18	Yes
199996	0.093259	No	22	Yes
199997	0.158304	No	19	Yes
199998	0.110972	No	24	Yes
199999	0.121960	No	22	Yes

[200000 rows x 12 columns]

```
[10]: #Scaling Numerical Features
from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()
data[['Claim Amount', 'Premium/Amount Ratio', 'BMI']] = scaler.
    ↪fit_transform(data[['Claim Amount', 'Premium/Amount Ratio', 'BMI']])
```

```
[11]: # Data Splitting
import pandas as pd
from sklearn.model_selection import train_test_split
# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(data, data['Churn'],
    ↪test_size=0.2, random_state=42)
# Print the shape of the training and testing sets
print("Training set shape:", X_train.shape)
print("Testing set shape:", X_test.shape)
```

Training set shape: (160000, 12)

Testing set shape: (40000, 12)

```
[12]: import pandas as pd
from sklearn.preprocessing import MinMaxScaler
# Load your dataset

data = pd.read_csv('randomdata.csv')
# Identify numerical columns in your dataset
numerical_columns = ['Claim Amount', 'Premium/Amount Ratio', 'BMI']
# Create a MinMaxScaler object
scaler = MinMaxScaler()
# Fit the MinMaxScaler to the numerical columns
scaler.fit(data[numerical_columns])
# Transform the numerical columns using the MinMaxScaler
data[numerical_columns] = scaler.transform(data[numerical_columns])
```

```
# Display the resulting DataFrame
print(data)
```

```

      Unnamed: 0    Customer Name \
0              0  Christine Payne
1              1    Tony Fernandez
2              2  Christopher Kim
3              3    Nicole Allen
4              4    Linda Cruz

```

```

...
199995      199995  Matthew Estrada
199996      199996      James Bean
199997      199997    David Meyer
199998      199998    Martha Stone
199999      199999    Shannon Lewis

```

```

      Customer_Address \
0    7627 Anderson Rest Apt. 265,Lake Heather, DC 3...
1    3953 Cindy Brook Apt. 147,East Lindatown, TN 4...
2          8693 Walters Mountains,South Tony, TX 88407
3          56926 Webster Coves,Shawnmouth, NV 04853
4    489 Thomas Forges Apt. 305,Jesseton, GA 36765

```

```

...
199995      2024 Lopez Gateway,Lake Pamelaafort, MS 35772
199996      0268 Lori Falls,West Jeffrey, SC 49142
199997      00573 Miller Cliff,New Allenbury, SC 68902
199998      62681 Peters Cove,South Anthony, RI 99783
199999      Unit 6569 Box 2236,DPO AE 88045

```

```

      Company Name Claim Reason Data confidentiality \
0  Williams, Henderson and Perez      Travel      Low
1          Moore-Goodwin      Medical      High
2          Smith-Holmes      Phone      Medium
3          Harrell-Perez      Phone      Medium
4  Simpson, Kramer and Hughes      Phone      Medium

```

```

...
199995      Carlson-Matthews      Medical      High
199996      Trevino-Cardenas      Medical      High
199997      Simon-Evans      Medical      High
199998      Baker, Brooks and Porter      Travel      Low
199999      Roth, Merritt and Grant      Medical      High

```

```

      Claim Amount  Category Premium  Premium/Amount Ratio \
0      0.163621      4794      0.309973
1      0.626197      14390      0.397222
2      0.110966      1875      0.545682
3      0.100957      1875      0.495739

```

4	0.103568	1875	0.508767
...
199995	0.679721	14390	0.432023
199996	0.583551	14390	0.369494
199997	0.990862	14390	0.634321
199998	0.231070	4794	0.441611
199999	0.763272	14390	0.486346

	Claim Request	output	BMI	Churn
0		No	0.3	Yes
1		No	0.6	Yes
2		No	0.0	Yes
3		No	0.6	Yes
4		No	0.3	Yes
...	
199995		No	0.0	Yes
199996		No	0.4	Yes
199997		No	0.1	Yes
199998		No	0.6	Yes
199999		No	0.4	Yes

[200000 rows x 12 columns]

DATA SPLITTING 80 20

```
[13]: import pandas as pd
from sklearn.model_selection import train_test_split

# Load your dataset

data = pd.read_csv('randomdata.csv')

# Assuming 'Churn' is your target variable and you want to predict it
X = data.drop(columns=['Churn']) # Features
y = data['Churn'] # Target variable

# Split the data into training and testing sets (80% training, 20% testing)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳ random_state=42)

# Display the sizes of the resulting sets
print(f"X_train shape: {X_train.shape}")
print(f"X_test shape: {X_test.shape}")
print(f"y_train shape: {y_train.shape}")
print(f"y_test shape: {y_test.shape}")
```

X_train shape: (160000, 11)

X_test shape: (40000, 11)

```
y_train shape: (160000,)
y_test shape: (40000,)
```

DATA SPLITTING 75 25

```
[14]: import pandas as pd
from sklearn.model_selection import train_test_split

# Load your dataset

data = pd.read_csv('randomdata.csv')

# Assuming 'Churn' is your target variable and you want to predict it
X = data.drop(columns=['Churn']) # Features
y = data['Churn'] # Target variable

# Split the data into training and testing sets (75% training, 25% testing)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25,
    ↪random_state=42)

# Display the sizes of the resulting sets
print(f"X_train shape: {X_train.shape}")
print(f"X_test shape: {X_test.shape}")
print(f"y_train shape: {y_train.shape}")
print(f"y_test shape: {y_test.shape}")
```

```
X_train shape: (150000, 11)
X_test shape: (50000, 11)
y_train shape: (150000,)
y_test shape: (50000,)
```

DATA SPLITTING 85 15

```
[15]: import pandas as pd
from sklearn.model_selection import train_test_split

# Load your dataset
# Replace 'randomdata.csv' with the actual path to your dataset
data = pd.read_csv('randomdata.csv')

# Assuming 'Churn' is your target variable and you want to predict it
X = data.drop(columns=['Churn']) # Features
y = data['Churn'] # Target variable

# Split the data into training and testing sets (85% training, 15% testing)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.15,
    ↪random_state=42)

# Display the sizes of the resulting sets
```



```

print(f"X_train shape: {X_train.shape}")
print(f"X_test shape: {X_test.shape}")
print(f"y_train shape: {y_train.shape}")
print(f"y_test shape: {y_test.shape}")

```

```

X_train shape: (170000, 11)
X_test shape: (30000, 11)
y_train shape: (170000,)
y_test shape: (30000,)

```

4 FEATURE SELECTION - Ensemble method

```

[16]: import pandas as pd
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder

data = pd.read_csv('randomdata.csv')

# Assuming 'Churn' is your target variable and you want to predict it
# Drop non-numeric columns and target variable
X = data.drop(columns=['Churn', 'Customer Name', 'Customer_Address'])
y = data['Churn'] # Target variable

# Encode categorical variables using LabelEncoder
label_encoder = LabelEncoder()
for column in X.columns:
    if X[column].dtype == 'object':
        X[column] = label_encoder.fit_transform(X[column])

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    random_state=42)

# Initialize a Random Forest Classifier
rf_classifier = RandomForestClassifier(n_estimators=100, random_state=42)

# Fit the classifier to your data
rf_classifier.fit(X_train, y_train)

# Evaluate the classifier on the testing data
accuracy = rf_classifier.score(X_test, y_test)
print("Accuracy:", accuracy)

```

Accuracy: 1.0

```

[17]: import pandas as pd
from sklearn.model_selection import train_test_split
from imblearn.over_sampling import RandomOverSampler
from sklearn.preprocessing import LabelEncoder
from sklearn.neural_network import MLPClassifier
from sklearn.metrics import accuracy_score, classification_report

data = pd.read_csv('randomdata.csv')

# Assuming 'Churn' is your target variable and you want to predict it
# Drop non-numeric columns and target variable
X = data.drop(columns=['Churn', 'Customer Name', 'Customer_Address'])
y = data['Churn'] # Target variable

# Encode categorical variables using LabelEncoder
label_encoder = LabelEncoder()
for column in X.columns:
    if X[column].dtype == 'object':
        X[column] = label_encoder.fit_transform(X[column])

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    random_state=42)

# Perform oversampling on the training data
oversampler = RandomOverSampler(sampling_strategy='auto', random_state=42)
X_train_resampled, y_train_resampled = oversampler.fit_resample(X_train,
    y_train)

# Initialize a Deep Learning model (MLPClassifier is used here as an example)
clf = MLPClassifier(hidden_layer_sizes=(100, 50), max_iter=1000,
    random_state=42)

# Fit the classifier to your resampled training data
clf.fit(X_train_resampled, y_train_resampled)

# Make predictions on the test data
y_pred = clf.predict(X_test)

# Evaluate the classifier
accuracy = accuracy_score(y_test, y_pred)
classification_report_result = classification_report(y_test, y_pred)

print("Accuracy:", accuracy)
print("Classification Report:")
print(classification_report_result)

```

Accuracy: 0.6986

Classification Report:

	precision	recall	f1-score	support
No	0.72	0.28	0.41	14536
Yes	0.70	0.94	0.80	25464
accuracy			0.70	40000
macro avg	0.71	0.61	0.60	40000
weighted avg	0.70	0.70	0.66	40000

5 DEEP LEARNING MODEL - Ensemble method

```
[18]: # Define a function to create a deep learning model
def create_model(input_dim):
    model = Sequential()
    model.add(Dense(64, input_dim=input_dim, activation='relu'))
    model.add(Dense(32, activation='relu'))
    model.add(Dense(1, activation='sigmoid'))
    model.compile(loss='binary_crossentropy', optimizer='adam',
metrics=['accuracy'])
    return model
```

```
[19]: # Create three deep learning models
import numpy as np
from tensorflow.keras.models import Sequential # Import the Sequential class
from tensorflow.keras.layers import Dense
from sklearn.metrics import accuracy_score
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler

model1 = create_model(X_train.shape[1])
model2 = create_model(X_train.shape[1])
model3 = create_model(X_train.shape[1])
```

```
[20]: from sklearn.preprocessing import LabelEncoder

# Encode the target variable into numerical values (0 and 1)
label_encoder = LabelEncoder()
y_train = label_encoder.fit_transform(y_train)
y_test = label_encoder.transform(y_test)

# Train the deep learning models
model1.fit(X_train, y_train, epochs=10, batch_size=32, verbose=0)
```

```
model2.fit(X_train, y_train, epochs=10, batch_size=32, verbose=0)
model3.fit(X_train, y_train, epochs=10, batch_size=32, verbose=0)
```

[20]: <keras.src.callbacks.History at 0x200bebbd6d0>

[21]: *# Make predictions using the individual models*

```
pred1 = model1.predict(X_test)
pred2 = model2.predict(X_test)
pred3 = model3.predict(X_test)
```

1250/1250 [=====] - 1s 621us/step

1250/1250 [=====] - 1s 601us/step

1250/1250 [=====] - 1s 610us/step

[22]: *# Ensemble the predictions using a simple averaging method*

```
ensemble_preds = np.round((pred1 + pred2 + pred3) / 3)
```

[23]: *# Calculate accuracy of the ensemble model*

```
ensemble_accuracy = accuracy_score(y_test, ensemble_preds)
```

```
print("Ensemble Model Accuracy:", ensemble_accuracy)
```

Ensemble Model Accuracy: 0.69195

[24]: `from sklearn.metrics import accuracy_score, precision_score, f1_score,
↪ recall_score`

Make predictions using the individual models

```
pred1 = model1.predict(X_test)
pred2 = model2.predict(X_test)
pred3 = model3.predict(X_test)
```

Ensemble the predictions using Voting Classifier

```
ensemble_preds = np.round((pred1 + pred2 + pred3) / 3)
```

Calculate classification metrics

```
ensemble_accuracy = accuracy_score(y_test, ensemble_preds)
ensemble_precision = precision_score(y_test, ensemble_preds)
ensemble_f1 = f1_score(y_test, ensemble_preds)
ensemble_recall = recall_score(y_test, ensemble_preds)
```

Print the evaluation metrics

```
print("Ensemble Model Metrics:")
print("Accuracy:", ensemble_accuracy)
print("Precision:", ensemble_precision)
print("F1 Score:", ensemble_f1)
print("Recall:", ensemble_recall)
```

```
1250/1250 [=====] - 1s 617us/step
1250/1250 [=====] - 1s 635us/step
1250/1250 [=====] - 1s 732us/step
Ensemble Model Metrics:
Accuracy: 0.69195
Precision: 0.6776137960860633
F1 Score: 0.8027217419148255
Recall: 0.984487904492617
```

Experimenting with another ensemble method - bagging

```
[25]: from sklearn.ensemble import BaggingClassifier, RandomForestClassifier

# Create a BaggingClassifier with a RandomForest base estimator
bagging_classifier = BaggingClassifier(
    base_estimator=RandomForestClassifier(n_estimators=100, random_state=42),
    n_estimators=10, random_state=42)

# Fit the bagging ensemble to the data
bagging_classifier.fit(X_train, y_train)

# Evaluate the bagging ensemble
bagging_accuracy = bagging_classifier.score(X_test, y_test)

print("Bagging Ensemble Accuracy:", bagging_accuracy)
```

```
C:\Users\Neha\anaconda3\Lib\site-packages\sklearn\ensemble\_base.py:166:
FutureWarning: `base_estimator` was renamed to `estimator` in version 1.2 and
will be removed in 1.4.
  warnings.warn(

Bagging Ensemble Accuracy: 1.0
```